

# End-to end Tracing, tracing Kubernetes

Aditya Chechani  
Joshua Manning  
Reet Chowdhary  
Runzhou Han

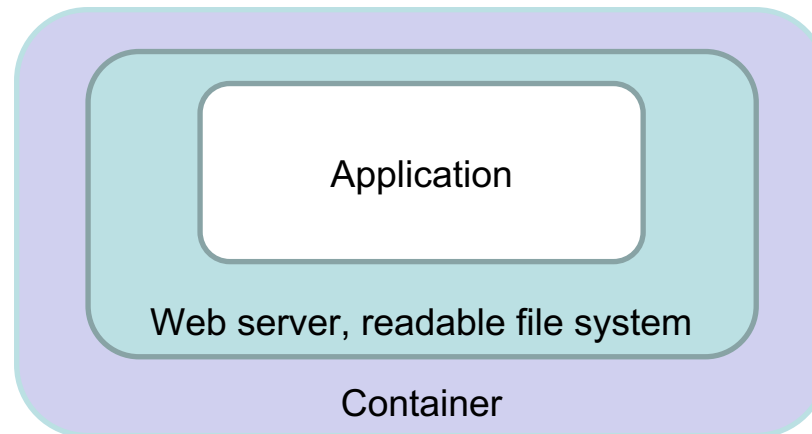
2/8/2018

Taiga: <https://tree.taiga.io/project/msdisme-2018-bucs528-template-7/>

GitHub: <https://github.com/BU-NU-CLOUD-SP18/Cross-Layer-Tracing-in-Kubernetes>

# Container

- Containers provide an isolated environment with a unique namespace
  - The environment satisfies a description of a set of resources required by an app
- Each App runs in a different container
  - An app can only use the resources defined in this namespace
  - App doesn't know what is outside its container



# Kubernetes

- Kubernetes is an orchestration system coordinating a highly available cluster of computers for deploying, scaling and managing containerized components of a distributed application in a datacenter. It makes these applications agnostic and isolated from other containers deployed on the same node e.g. machine.
  - It allocates resources to containers to have them meet the required description
  - It offers a unique namespace to each container
  - It can scale in & scale out an application and make replication) when required
  - Provides an abstraction through which each container is able to communicate with the outside world

# But why Kubernetes?


- Better management through modularity.
  - Faster Development.
  - Achieved using Pods.
- Deploying and updating software at scale.
  - Scalability
  - Visibility
  - Time Saving
  - Version Control

# Components of Kubernetes

- Containers
- Control Plane
  - The decision making part of Kubernetes which decides what has to be done with each containers according to its description
- Data Plane
  - This is the part which enforces all of control plane's decisions
  - Allows applications to remain agnostic to their surroundings
- e.g: Post Service

# Tracing

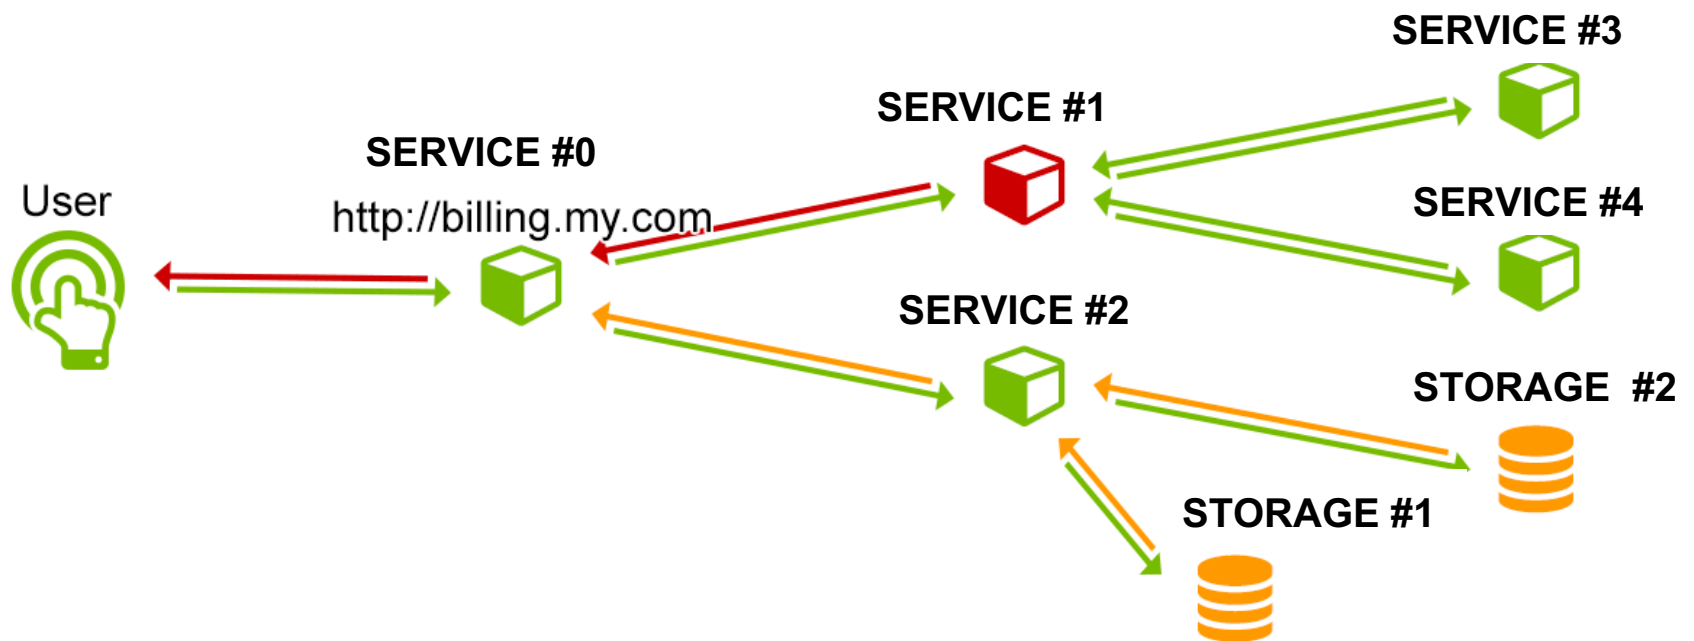
- To understand the performance/correctness of deploying distributed applications, many parts e.g application, networks, dataplane, control plane etc have to be understood.
- Tracing: Trace events labeled by a unique ID of a request to know apps' behavior
  - Tracing ID is a local variable generated by an app
  - Propagation of request IDs tied to a specific request



# Tracing

- Three major components of Tracing
  - Trace points
  - Tracing ID propagation
  - Huge repository storage to store tracing information







# End to End Tracing

- Tracing is determining the behaviour of each application/container. But what if it is not the application/container which has a latency issue?
- End to End (e2e) Tracing is to follow the execution of requests infrastructure along the entirety of its “**PATH**” of propagation (including its dataplane), to provide detailed tracing for capacity planning and performance analysis

# Vision & Goals

- End-to-end tracing for app behavior on Kubernetes
  - Trace points (start time, stop time, annotations for remote procedure call)
  - RPC IDs
- Identify issues (bottle necks, latency issues etc.)
- Add in trace points to “data plane” features
  - Ingress - An API object that manages external access to the services in a cluster, typically HTTP.
  - Services - An abstraction which defines a logical set of **Pods** and a policy by which to access them - sometimes called a micro-service.

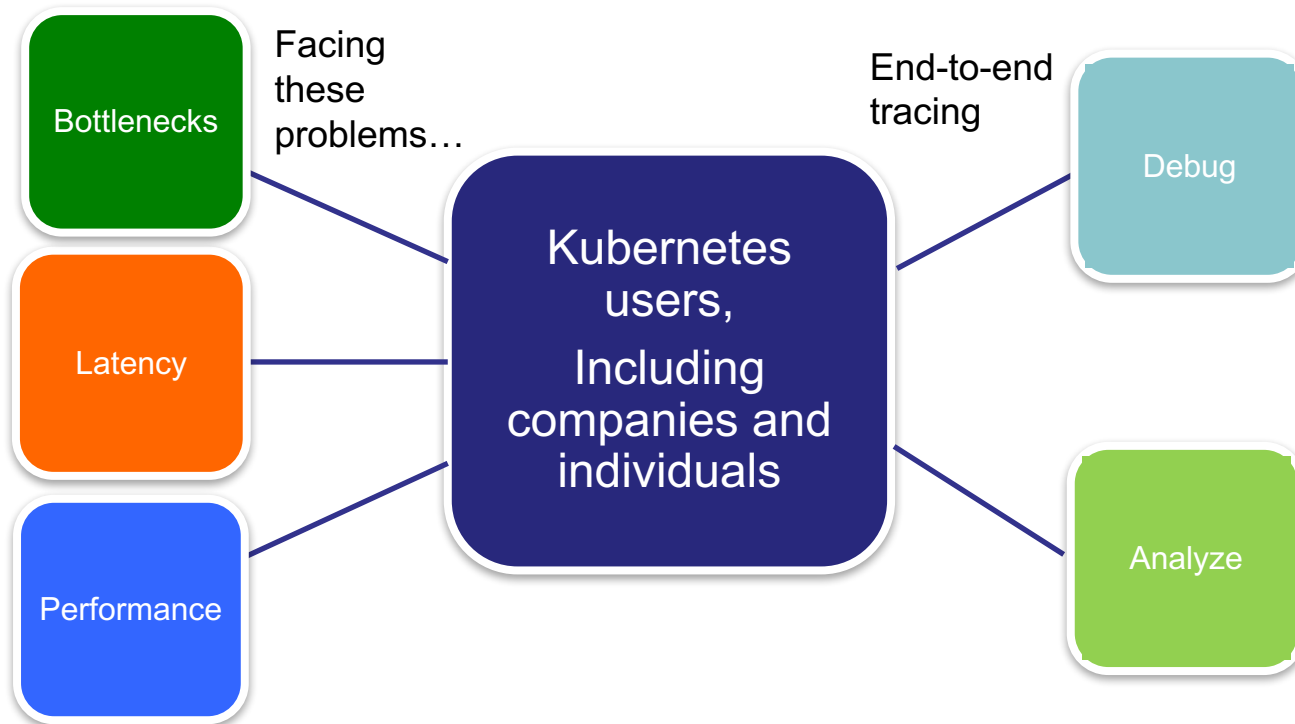
# Scope & Features

- Instrument the services section of Kubernetes with trace points
- Instrument the ingress features of Kubernetes with trace points
- Focus on data-plane portion

# User Stories

- A developer would be able to debug and analyze their deployed distributed applications
- Kubernetes developers can look into which requests are present in a particular data plane at some time  $T$ .
- Performance of a node can be determined
- Cloud providers and data centers

# Our Users(This part is not sure)



# Challenges

- Tracing ID is from an app inside a container, how to let it propagate between containers and nodes?
- How to make tracing events visible and record them?

# Example

- Jaeger: a distributed tracing system used for monitoring microservices-based distributed systems donated by Uber.
- [OpenTracing](#): A new, open standard for instrumenting applications and OSS packages for distributed tracing and monitoring.
- HotRod: a sample application that can find out drivers nearby.

# Example

Your web client's id: 3320

A random ID

Client

Hot R.O.D.  
*Rides On Demand*

Rachel's Floral Designs

Trom Chocolatier

Japanese Deserts

Amazing Coffee Roasters

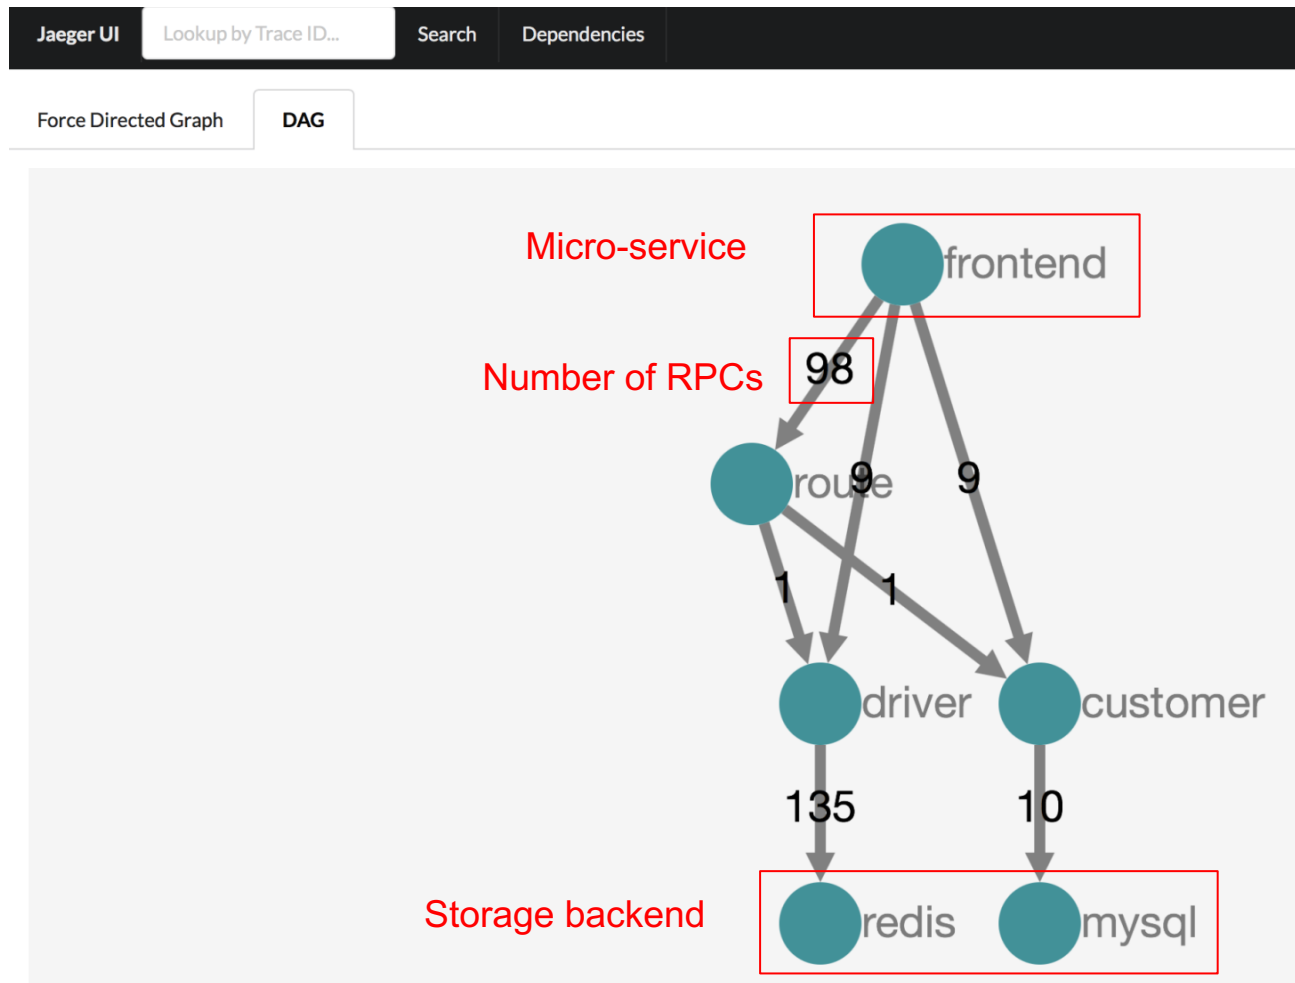
Click on customer name above to order a car.

Tracing ID

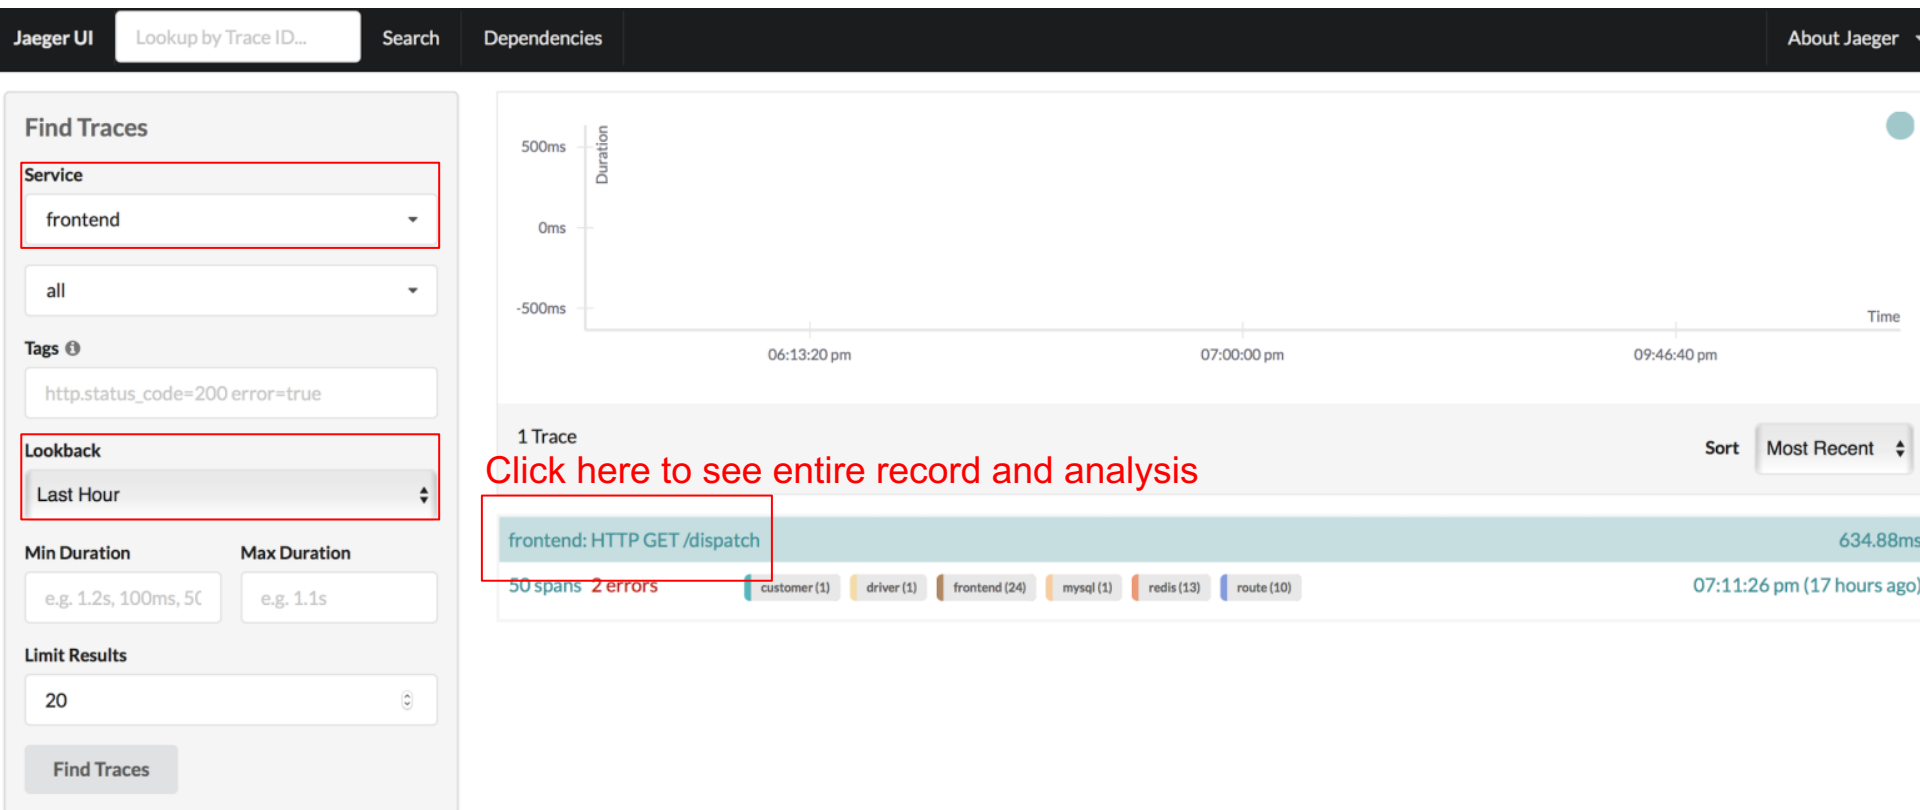
HotROD **T743182C** arriving in 2min [req: 3320-10, latency: 640ms]  
HotROD **T707383C** arriving in 2min [req: 3320-9, latency: 767ms]  
HotROD **T749520C** arriving in 2min [req: 3320-8, latency: 734ms]  
HotROD **T702496C** arriving in 2min [req: 3320-7, latency: 784ms]  
HotROD **T737493C** arriving in 2min [req: 3320-6, latency: 783ms]  
HotROD **T782215C** arriving in 2min [req: 3320-5, latency: 740ms]  
HotROD **T737650C** arriving in 3min [req: 3320-4, latency: 708ms]  
HotROD **T783781C** arriving in 2min [req: 3320-3, latency: 762ms]  
HotROD **T758052C** arriving in 2min [req: 3320-2, latency: 733ms]  
HotROD **T791209C** arriving in 2min [req: 3320-1, latency: 872ms]



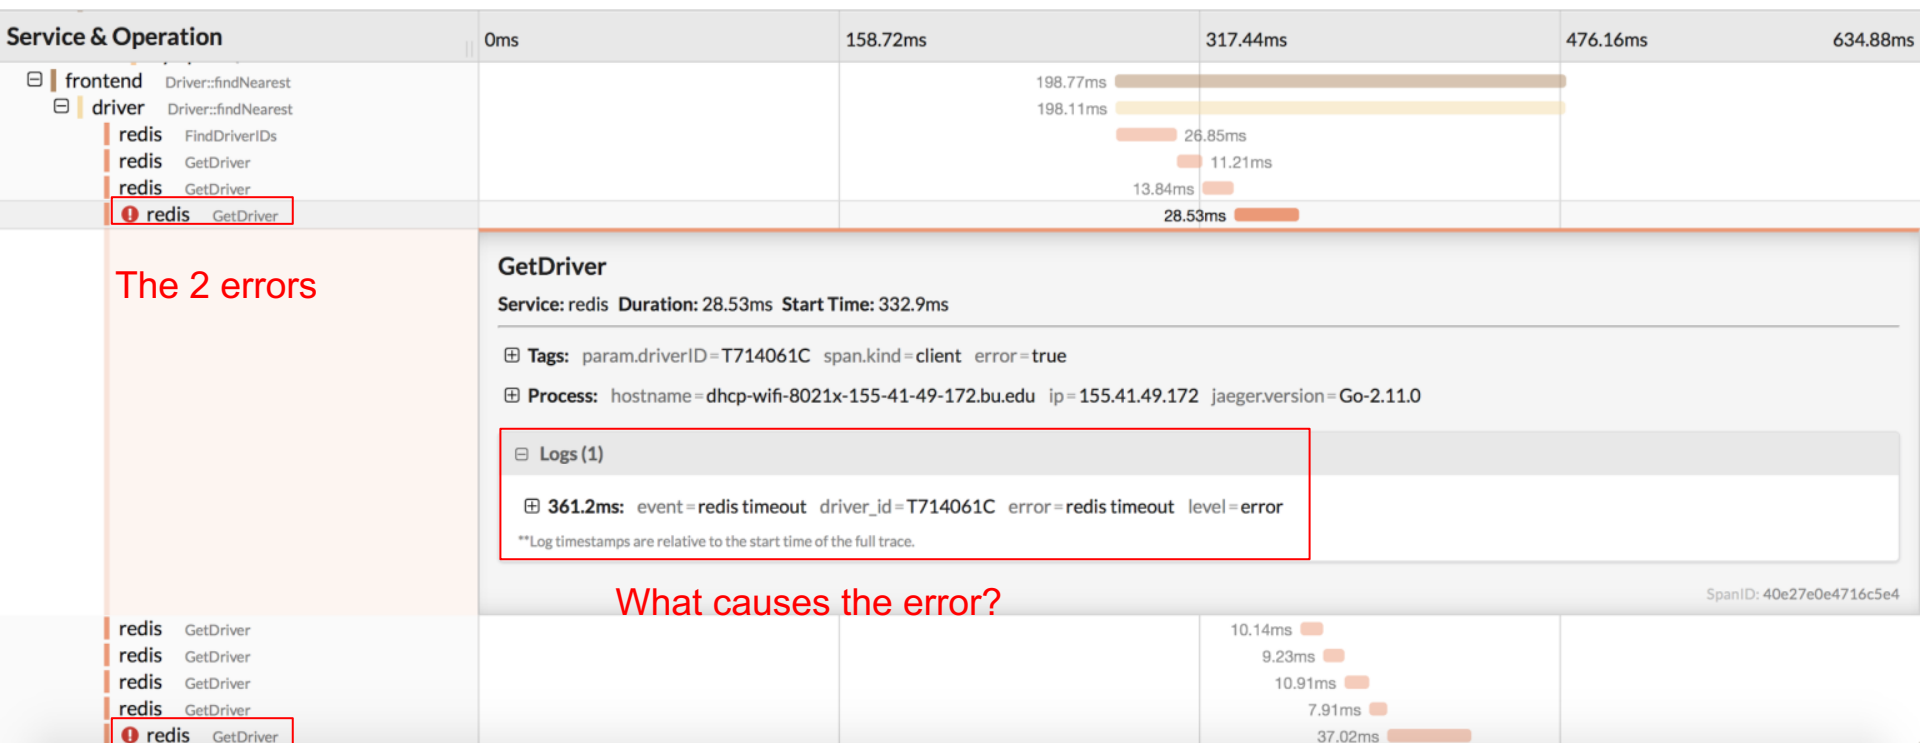
# Example



# Example



# Example



# Example

Jaeger can detect latency and bottlenecks and find out how to deal with them.

Let's go to the HotROD UI and click on one of the buttons repeatedly (and quickly).we can see, the more requests are being processed concurrently, the longer it takes for the backend to respond.

Your web client's id: 3320

## Hot R.O.D.

*Rides On Demand*

Rachel's Floral Designs

Trom Chocolatier

Japanese Deserts

Amazing Coffee Roasters

Click on customer name above to order a car.

HotROD **T769232C** arriving in 2min [req: 3320-40, latency: 3945ms]  
HotROD **T776979C** arriving in 2min [req: 3320-39, latency: 3882ms]  
HotROD **T773478C** arriving in 2min [req: 3320-38, latency: 3661ms]  
HotROD **T790373C** arriving in 2min [req: 3320-37, latency: 3623ms]  
HotROD **T767379C** arriving in 2min [req: 3320-36, latency: 3373ms]  
HotROD **T715828C** arriving in 2min [req: 3320-35, latency: 3267ms]  
HotROD **T760959C** arriving in 2min [req: 3320-34, latency: 2976ms]  
HotROD **T796518C** arriving in 2min [req: 3320-33, latency: 2885ms]  
HotROD **T761333C** arriving in 2min [req: 3320-32, latency: 2882ms]

More information click [here](#)

# Acceptance Criteria

Propagating a request through the data plane of Kubernetes while making Kubernetes knowledgeable about the metadata being propagated along with the request.

Stretching the capabilities of the already existing tracing systems to the network namespace layer of Kubernetes which could prove to cause a substantial increase in latency and reduction in performance.

# MVP(minimum value product)

- Adding trace points to Kubernetes on TWO data paths of a distributed application in parallel and then have them show up in Jaeger.

# Release Planning

## Short Term Goals:

- Get Kubernetes on our machines
- Tutorials on Kubernetes & tracing
- Containerize Hotrod using a container providing platform. e.g docker.
- Implement Jaeger in this containerized Hotrod application as backend to look up the traces and trace points.
- Study up on Kubernetes dataplane subsystems e.g HAProxy.

**THANK YOU**