

# CROSS LAYER TRACING IN KUBERNETES

Runzhou Han

Aditya Chechani

Reet Chowdhary

Taiga: <https://tree.taiga.io/project/msdisme-2018-bucs528-template-7/>

GitHub: <https://github.com/BU-NU-CLOUD-SP18/Cross-Layer-Tracing-in-Kubernetes>


# Review

- Kubernetes
  - Ingress
    - Ingress controller
      - Nginx
      - Haproxy
  - Service
- Tracing
  - Data plane

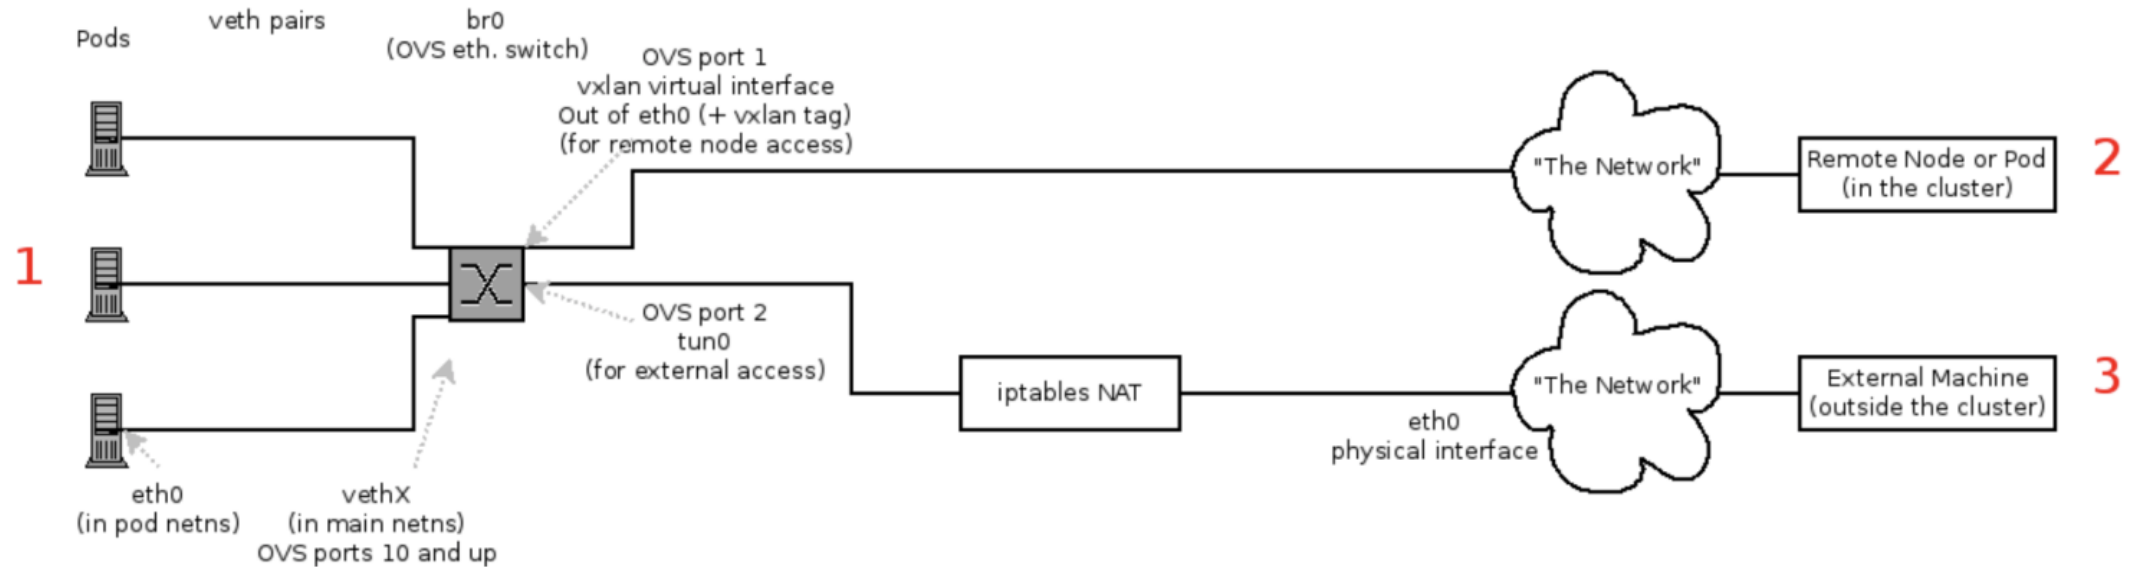
# Task Completed

1. Setup MOC
2. Deploy Kubernetes
3. Trace HotRod using Jaeger on MOC and locally
4. Get logs from Kubernetes Control Plane and make sense of it
5. Create a flask application and deploy in Kubernetes
6. Get Nginx Logs locally

# Kubernetes Data Plane



# Kubernetes Data Plane



# Ingress

- Ingress is a collection of rules that allow inbound connections to reach the internal cluster services
- In order for ingress to work the cluster must have an Ingress controller running
- Currently GCE and nginx are supported, some instances of the documentation indicates that HAproxy is the load balancer
- The plan is to add trace points to nginx or HAproxy in the next sprint in order to trace ingress

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: test
  annotations:
    ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
    - host: foo.bar.com
      http:
        paths:
          - path: /foo
            backend:
              serviceName: s1
              servicePort: 80
          - path: /bar
            backend:
              serviceName: s2
              servicePort: 80
```

When you create the Ingress with `kubectl create -f :`

```
$ kubectl get ing
NAME      RULE      BACKEND  ADDRESS
test      -         foo.bar.com
          /foo      s1:80
          /bar      s2:80
```



# Containerizing and Deploying Flask

- In order to display the functionality of Kubernetes Ingress features we decided to use a Flask app
- The application returns a simple HTML response upon receiving a GET request
- We were able to containerize the application and deploy it and receive a response from it on the local network
- We must now expose the application to the outside world and send the requests through the ingress

# Change in Plan:

Deploy a simple containerized flask application in Kubernetes

Establish communication from outside world

Get logs/traces from the Ingress Controller (Nginx).

```
# nginx.conf
http {
    ...
    log_format combined '$remote_addr - $remote_user [$time_local] '
                        '$request' $status $body_bytes_sent '
                        '$http_referer' '$http_user_agent';
    ...
}
```

Figure out where the trace points can be added in Nginx

If time permits, look into OVS switch which is responsible for communication between nodes inside Kubernetes.

# Some Issues

## Frequent Connection issues

```
Reetc@DESKTOP-KJMHN2I /cygdrive/c/Users/Reetc/Desktop  
$ ssh -A centos@128.31.26.26  
ssh: connect to host 128.31.26.26 port 22: Connection timed out
```

Recently kubectl is timing out after we're already on the master node

```
^C[centos@192 ~]$ kubectl get nodes  
Unable to connect to the server: dial tcp 128.31.24.168:8080: i/o timeout  
[centos@192 ~]$ Connection reset by 128.31.26.26 port 22
```

# Taiga Burndown chart




76%

100 project points

100 defined points

76 closed points

8 points / sprint



# Release Planning

■ In sprint 5, we will:

- Add tracing points to Kubernetes ingress controller
- Show that we can collect tracing events of Kubernetes

■ In final sprint, we will:

- Compare Kubernetes tracing and application tracing

# Questions?